

# Bayesian inference with **JAGS** and `rjags`

M. L. Delignette-Muller  
VetAgro Sup - LBBE

September 15, 2021



# Clostridium example

## Modeling of the dose-response curve related to the ingestion of *Clostridium perfringens*.

- **Deterministic part** of the model, probability that the host gets sick:

$$p = 1 - (1 - r)^{dose}$$

with *dose* le number of ingested cells

- **Stochastic part** of the model, number of sick hosts *N<sub>sick</sub>* for *N* exposed hosts :

$$N_{sick} \sim \text{Binomiale}(n = N, p = 1 - (1 - r)^{dose})$$

# Formalization of a model using a DAG - Directed Acyclic Graph

## What is a DAG ?

- **a directed graph**  
*(all the links are directed)*
- **without cycles (loops)**  
*(from each node, and following the links, it is impossible to return to this node)*
- that we use in Bayesian inference to represent **conditional dependencies** between nodes.  
*(you can see a DAG as a mecanistic description of how output data could be used simulated from input data.)*

# DAG formalism

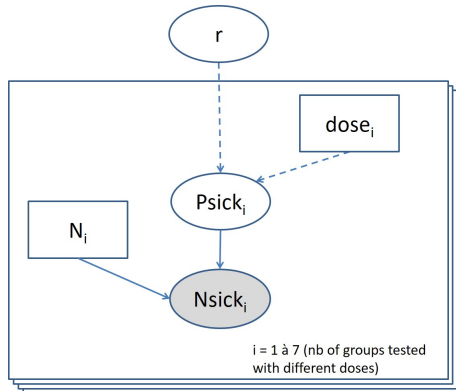
## ■ Nodes

- covariable (rectangle)
- variable (ellipse)
  - observed variable, latent variable or intermediate variable
  - Variables corresponding to output data are sometimes shaded

## ■ Links

- deterministic link (or logical link - dashed arrow - link that could be omitted by writing the model more synthetically)
- stochastic link (solid line arrow - essential link, that cannot be omitted)

# DAG of the model on our example



Mathematical definition of links

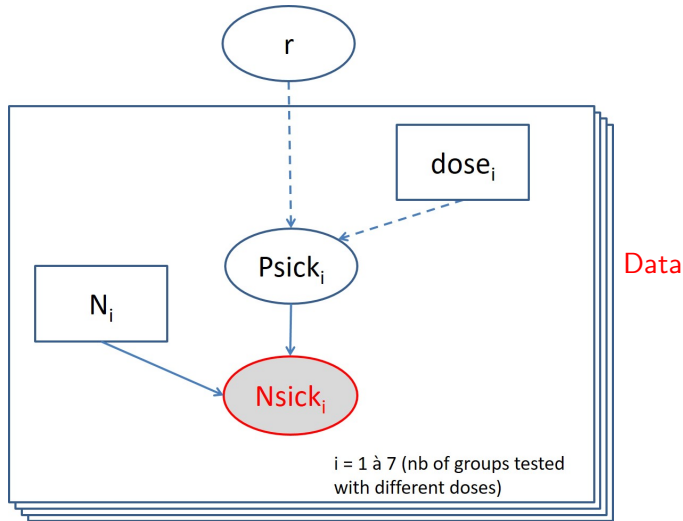
- Deterministic links

$$Psick_i = 1 - (1 - r)^{dose_i}$$

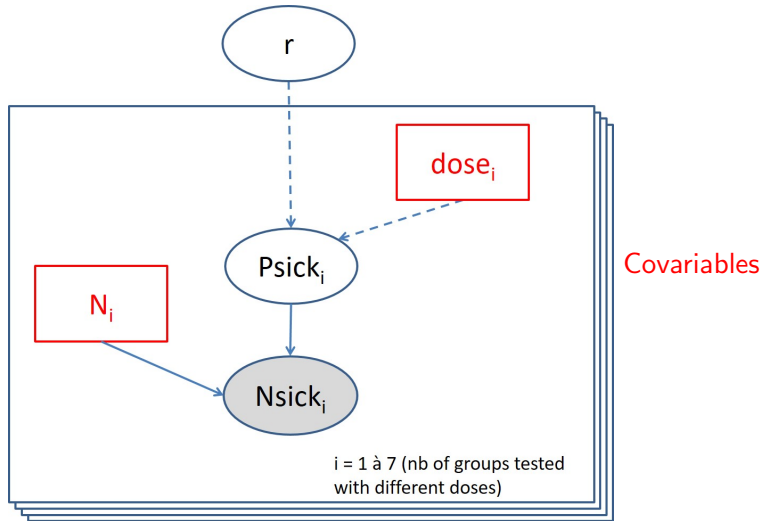
- Stochastic links

$$Nsick_i \sim Binomiale(N, Psick_i)$$

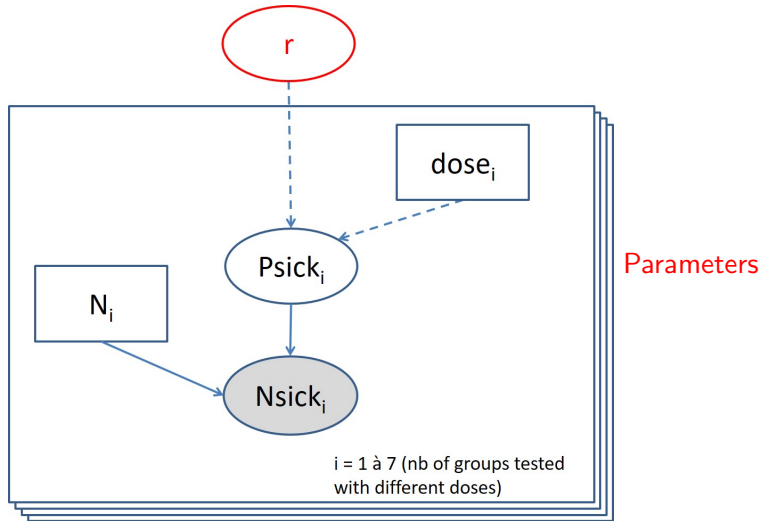
# DAG of the model - data (likelihood)



# DAG of the model - covariables (explicative variables)



# DAG of the model - parameters (to estimate)





# Prior information


In this example, we will assume it is reasonable to define from prior information about the unique parameter:

- a uniform prior distribution between -15 and -5 on  $\log_{10}(r)$ ,

# Data related to our example

Number of sick persons  $N_{sick_i}$  for each group of  $N_i$  persons  
exposed at the dose  $dose_i$

```
> plot( $N_{sick}/N \sim dose_{log10}$ , data = d, pch = 16,  
+ xlab = " $log_{10}(dose)$ ", ylab = " $proportion\ of\ sick\ persons$ ")
```



# The BUGS project (since 1989)


## Bayesian inference Using Gibbs Sampling

Development and provision of flexible software to implement Bayesian inference on complex models using MCMC.

Some available tools :

- WinBUGS and OpenBUGS
- **JAGS (Just Another Gibbs sampler - Martyn Plummer)**
- Stan and Nimble (new algorithms added to MCMC)
- RevBayes (for phylogeny)
- several other tools for specific model families

# Evolution of the number of PubMed citations with **Bayesian** in the title from the beginning of the project



# Coding of a model in the BUGS language

## A declarative language

(the order of the command lines does not matter)  
that looks like **R**

### ■ Declaration of a deterministic node

```
node <- fonction(some other nodes)
```

### ■ Declaration of a stochastic node

including input nodes,

i.e. parameters stochastically defined by their prior

```
node ~ distribution(optionnally some other nodes)
```

**BE CAREFUL:** a node on which we have data must always be coded by a stochastic link !

## Code of the model in our example

To be written in a text file or in a string as below.

```
> model <-  
+ "model  
+ {  
+   for (i in 1:Ndose)  
+   {  
+     psick[i] <- 1 - (1 - r)^dose[i]  
+     Nsick[i] ~ dbin(psick[i], N[i])  
+   }  
+   log10r ~ dunif(-15, -5)  
+   r <- 10^log10r  
+ }  
+ "
```

# Some properties of the BUGS language that differentiate it from R

A node is univariate.

It is necessary to specify the dimensions, the indices, and **explicitly write loops** to define vectors or matrices or multidimensional arrays.

For example, we can write:

```
v[]      v[i]
M[, ]    M[i, j]
A[, , ,] A[i, j, k, l]
M[, j]   v[n:m]
x[y[i]]  x[2*j-1]
```

# Let us build the code of our model step by step

A loop to define all the observations

```
model
{
  for(i in 1:Ndose)
  {
    Nsick[i] ~ dbin(psick[i], N[i])
  }
}
```



## Build of the code - add of intermediate variables

All nodes must be defined in the model except covariables.  
The order of lines dose not matter.

```
model
{
  for(i in 1:Ndose)
  {
    Nsick[i] ~ dbin(psick[i], N[i])
    psick[i] <- 1 - (1 - r)^dose[i]
  }
}
```

## Build of the code - add of priors

Prior distributions of parameters (here just one) must be defined outside the loop.

```
model
{
  for(i in 1:Ndose)
  {
    Nsick[i] ~ dbin(psick[i], N[i])
    psick[i] <- 1 - (1 - r)^dose[i]
  }
  log10r ~ dunif(-15, -5)
  r <- 10^log10r
}
```

## Other differences between **BUGS** and **R** languages

BE CAREFUL,  
the BUGS language and the R language are different,  
and **some differences concern the name of the distributions  
and their parameterization.**

Refer to the user manual of JAGS or of other languages for a  
complete and up-to-date list of the functions and distributions.  
The JAGS reference manual:

http:

[//sourceforge.net/projects/mcmc-jags/files/Manuals/](http://sourceforge.net/projects/mcmc-jags/files/Manuals/)

## Coding of data

Coding of data is software-dependent.

Here we will use **JAGS** (MCMC) and **rjags**.

Data must be defined in a data list (here named `data4jags`).

```
> require(rjags)

> data4jags <- list(dose = 10^d$dosedlog10,
+                  N = d$N,
+                  Nsick = d$Nsick,
+                  Ndose = nrow(d))
```

# Pay attention to the consistency between the names used in the model and in the data list

- All the nodes appearing in the model but not defined in the model, so appearing only to the right of an operator, (here *dose* and *N*)
- as well as the max loop indices (here *Ndose*)
- and the output of the model (observed data, here *Nsick*)

must be defined in the data list.

**BE CAREFUL to use the same names in the data list and the model code !**

# Definition of MCMC initial values

Software-dependent coding.

(described here for **JAGS** and **rjags**)

The **definition of initial values** is theoretically required **for each input node and each chain** especially for a correct use of the **Gelman and Rubin statistics** to appreciate the convergence of MCMCs (otherwise, for each parameter, the chains all start from the same value defined by default as a central value of its prior distribution).

**Ex.**

```
> ini <- list(list(log10r = -12),  
+           list(log10r = -11),  
+           list(log10r = -10))
```

# Simulations

## ■ Build of a model and adaptation

```
> m <- jags.model(file = textConnection(model),  
+               data = data4jags, inits = ini,  
+               n.chains = 3, n.adapt = 1000)
```

`n.adapt` (fixed by default to 1000) corresponds to the number of iterations of a phase during which the algorithm is adapted, so during which the simulated values are not yet MCMCs.

## ■ Burnin phase

```
> update(m, 3000)
```

## ■ Monitoring of simulations

```
> mc <- coda.samples(m, c("r"), n.iter = 1000)  
> # generally one starts rather with n.iter around 5000
```

# Simulations

## ■ Build of a model and adaptation

```
> m <- jags.model(file = textConnection(model),  
+                 data = data4jags, inits = ini,  
+                 n.chains = 3, n.adapt = 1000)
```

n.adapt (fixed by default to 1000) corresponds to the number of iterations of a phase during which the algorithm is adapted, so during which the simulated values are not yet MCMCs.

## ■ Burnin phase

```
> update(m, 3000)
```

## ■ Monitoring of simulations

```
> mc <- coda.samples(m, c("r"), n.iter = 1000)  
> # generally one starts rather with n.iter around 5000
```



# Simulations

## ■ Build of a model and adaptation

```
> m <- jags.model(file = textConnection(model),  
+               data = data4jags, inits = ini,  
+               n.chains = 3, n.adapt = 1000)
```

`n.adapt` (fixed by default to 1000) corresponds to the number of iterations of a phase during which the algorithm is adapted, so during which the simulated values are not yet MCMCs.

## ■ Burnin phase

```
> update(m, 3000)
```


## ■ Monitoring of simulations

```
> mc <- coda.samples(m, c("r"), n.iter = 1000)  
> # generally one starts rather with n.iter around 5000
```

# MCMC trace

All chains must converge to the same limit in term of distribution (stability and overlap/good mixing of the chains).  
Here the mixing seems acceptable.

```
> plot(mc, density = FALSE)
```



# Gelman-Rubin convergence diagnostic

For each parameter, defined by the square root of the ratio between the variance of its posterior marginal distribution and the intra-chain variance, which we expect to be 1 when convergence is reached.


Gelman indicates 1.1 as a maximum acceptable value for all nodes while indicating that one should try to reach 1.00 to get precise final results from MCMCs.

```
> gelman.diag(mc)
```

Potential scale reduction factors:

	Point est.	Upper C.I.
r	1	1.01

# Example of MCMC chains with a bad overlap



```
> gelman.diag(mc3.3c)
```

Potential scale reduction factors:


	Point est.	Upper C.I.
110alpha	1.01	1.03

# Autocorrelation plot


For each chain, plot of the correlation between MCMC iterations as a function of the lag between iterations.

**Here the autocorrelation is very low.**


```
> autocorr.plot(mc[[1]])
```




# Trace a chain with an acceptable low autorrelation



# Trace of a chain with a stronger autocorrelation that would need a thinning




# Autocorrelation plot for this chain






# Principle of thinning

With a thin of 10 one stores 1 iteration out of 10.  
A thinned chain may contain most of the information when taking up less space in memory.




## Principle of thinning (2)

After thinning: 100 out of 1000 iterations.



## Principle of thinning (3)


After thinning the number of iterations is low (here only 100).



## Principle of thinning (4)


It is thus necessary to increase the initial number of iterations (here  $\times 10 \rightarrow$  longer computation).

```
> mc3.1c <- coda.samples(m3.1c, c("l10alpha"), n.iter = 10000, thin = 10)  
> plot(mc3.1c, density = FALSE, main = "")
```



# Visualisation of the posterior distribution

```
> plot(mc, trace = FALSE)
```



# Statistical summary

```
> summary(mc)
```

```
Iterations = 4001:5000
```

```
Thinning interval = 1
```

```
Number of chains = 3
```

```
Sample size per chain = 1000
```

1. Empirical mean and standard deviation for each variable,  
plus standard error of the mean:

Mean	SD	Naive SE	Time-series SE
4.96e-11	1.39e-11	2.53e-13	0.00e+00

2. Quantiles for each variable:

2.5%	25%	50%	75%	97.5%
2.69e-11	3.95e-11	4.81e-11	5.82e-11	8.00e-11

# Credibility intervals

## ■ Classically based on 2.5% and 97.5% quantiles

```
> summary(mc)$quantiles
```

2.5%	25%	50%	75%	97.5%
2.69e-11	3.95e-11	4.81e-11	5.82e-11	8.00e-11


## ■ Less classical High Posterior Density (HPD) intervals

```
> HPDinterval(mc[[1]], prob = 0.95) # here for the first chain
```

	lower	upper
r	2.58e-11	7.84e-11

```
attr("Probability")  
[1] 0.95
```

# Difference between both intervals for asymmetrical posterior distributions





# Conclusion

**Now it's your turn to play with JAGS !**

To learn the technical aspects, nothing is best than practice !



You have an introductory guide to **JAGS** and `rjags` to help you to start and go further in particular for prediction and model validation.